

Decision Tree

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
```

```
In [4]: df = pd.read_csv('data/data_no_fliers.csv')
df.drop('Unnamed: 0', axis=1, inplace=True)
df.head()
```

```
Out[4]:
```

	neo	pha	H	epoch	epoch_mjd	epoch_cal	e	a	q	i
0	0	0	3.40	2458600.5	58600	20190427.0	0.076009	2.769165	2.558684	10.594067
1	0	0	4.20	2459000.5	59000	20200531.0	0.229972	2.773841	2.135935	34.832932
2	0	0	5.33	2459000.5	59000	20200531.0	0.256936	2.668285	1.982706	12.991043
3	0	0	3.00	2458600.5	58600	20190427.0	0.088721	2.361418	2.151909	7.141771
4	0	0	6.90	2459000.5	59000	20200531.0	0.190913	2.574037	2.082619	5.367427

5 rows × 45 columns

```
In [4]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 921430 entries, 0 to 921429
Data columns (total 45 columns):
#   Column          Non-Null Count  Dtype
---  -
0   neo              921430 non-null  int64
1   pha              921430 non-null  int64
2   H                921430 non-null  float64
3   epoch            921430 non-null  float64
4   epoch_mjd        921430 non-null  int64
5   epoch_cal        921430 non-null  float64
6   e                921430 non-null  float64
7   a                921430 non-null  float64
8   q                921430 non-null  float64
9   i                921430 non-null  float64
10  om               921430 non-null  float64
11  w                921430 non-null  float64
12  ma               921430 non-null  float64
13  ad               921430 non-null  float64
14  n                921430 non-null  float64
15  tp               921430 non-null  float64
16  tp_cal           921430 non-null  float64
```

```

17 per          921430 non-null float64
18 per_y        921430 non-null float64
19 moid         921430 non-null float64
20 moid_ld      921430 non-null float64
21 sigma_e      921430 non-null float64
22 sigma_a      921430 non-null float64
23 sigma_q      921430 non-null float64
24 sigma_i      921430 non-null float64
25 sigma_om     921430 non-null float64
26 sigma_w      921430 non-null float64
27 sigma_ma     921430 non-null float64
28 sigma_ad     921430 non-null float64
29 sigma_n      921430 non-null float64
30 sigma_tp     921430 non-null float64
31 sigma_per    921430 non-null float64
32 rms          921430 non-null float64
33 class_AMO    921430 non-null int64
34 class_APO    921430 non-null int64
35 class_AST    921430 non-null int64
36 class_ATE    921430 non-null int64
37 class_CEN    921430 non-null int64
38 class_IEO    921430 non-null int64
39 class_IMB    921430 non-null int64
40 class_MBA    921430 non-null int64
41 class_MCA    921430 non-null int64
42 class_OMB    921430 non-null int64
43 class_TJN    921430 non-null int64
44 class_TNO    921430 non-null int64
dtypes: float64(30), int64(15)
memory usage: 316.3 MB

```

Iteration 1: Baseline

```

In [5]: X = df.drop(['pha'], axis=1)
        y = df['pha']

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random

In [6]: tree_pipe = Pipeline([('ss', StandardScaler()),
                              ('clf', DecisionTreeClassifier(random_state=123, class_weig

In [7]: tree_pipe.fit(X_train, y_train)
        y_pred = tree_pipe.predict(X_test)

In [8]: print('Train Report')
        print(classification_report(y_train, tree_pipe.predict(X_train)))
        print('\n')
        print('Test Report')
        print(classification_report(y_test, y_pred))

```

Train Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	1.00	1.00	1.00	1524
accuracy			1.00	691072
macro avg	1.00	1.00	1.00	691072
weighted avg	1.00	1.00	1.00	691072

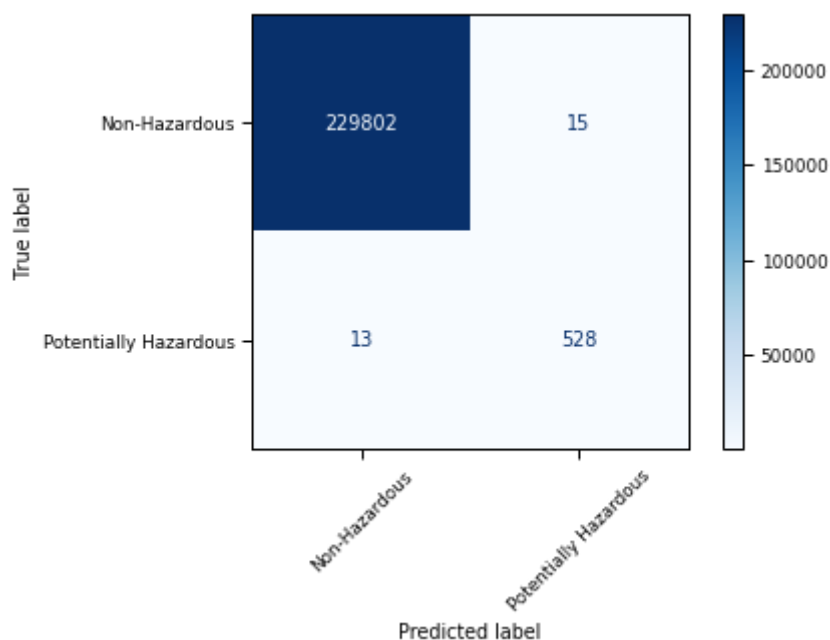
Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.97	0.98	0.97	541
accuracy			1.00	230358
macro avg	0.99	0.99	0.99	230358
weighted avg	1.00	1.00	1.00	230358

The testing metrics look promising, but the model is overfit to the training data. The baseline model performs better than Iteration 3 of logistic regression. The macro avg for recall is 1% lower than it was in logistic regression: Iteration 3, and recall for the positive class is 2% lower, but these values will likely go up with enough model tuning.

```
In [58]: sns.set_context('paper')

plot_confusion_matrix(tree_pipe, X_test, y_test, cmap=plt.cm.Blues,
                      display_labels=['Non-Hazardous', 'Potentially Hazardous'],
```



This model has fewer false positives than logistic regression, and I expect the number of false negatives to decrease as the model is tuned while prioritizing recall.

```
In [10]: tree_pipe.steps[1][1].get_params()
```

```
Out[10]: {'ccp_alpha': 0.0,
          'class_weight': {0: 1, 1: 200},
          'criterion': 'gini',
          'max_depth': None,
          'max_features': None,
          'max_leaf_nodes': None,
          'min_impurity_decrease': 0.0,
          'min_impurity_split': None,
          'min_samples_leaf': 1,
          'min_samples_split': 2,
          'min_weight_fraction_leaf': 0.0,
          'presort': 'deprecated',
          'random_state': 123,
          'splitter': 'best'}
```

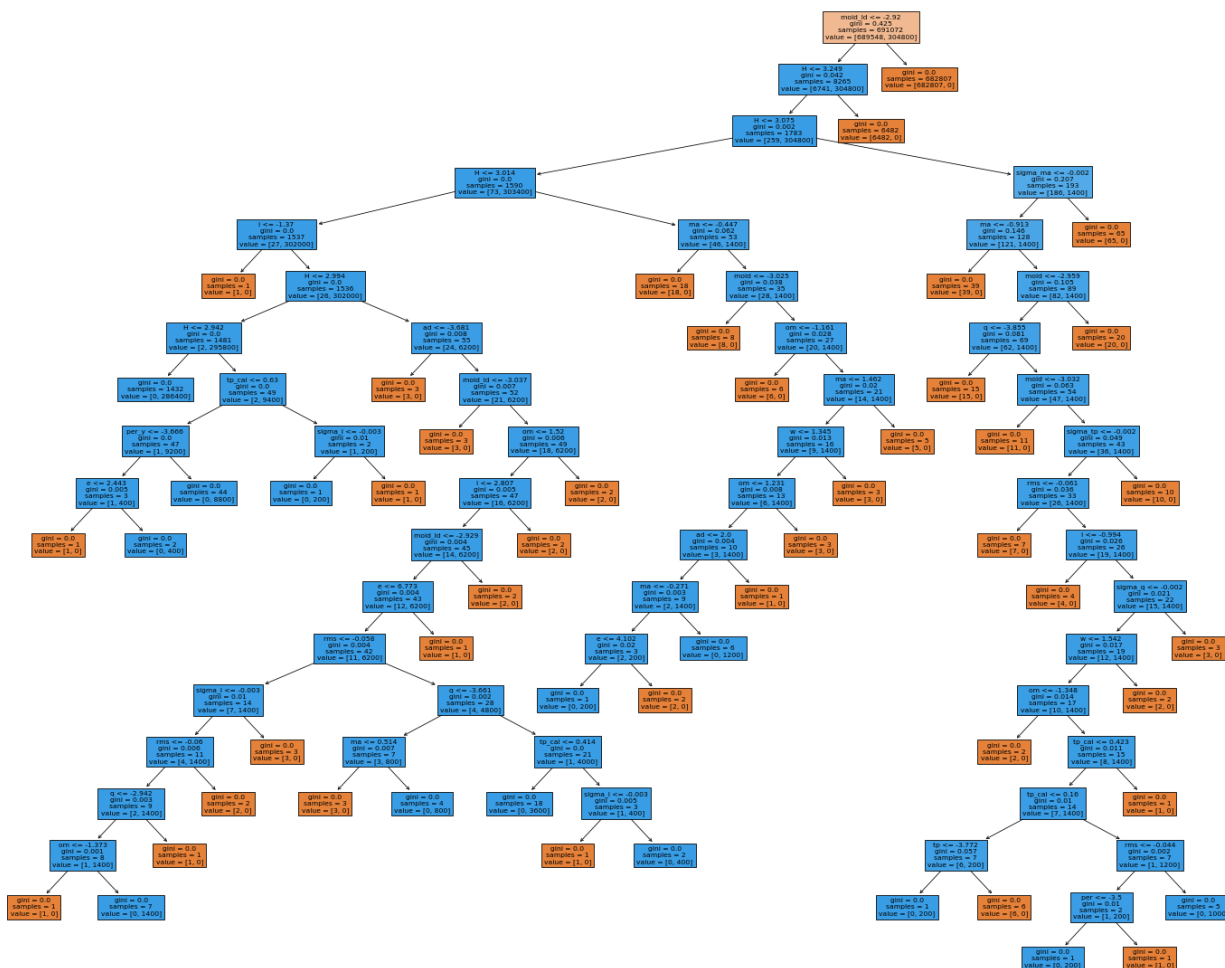
```
In [11]: tree_pipe.steps[1][1].get_depth()
```

```
Out[11]: 18
```

The initial tree has a max depth of 18. Will try some lower values in the next iteration in order to address overfitting.

```
In [12]: from sklearn.tree import plot_tree
```

```
In [59]: fig= plt.figure(figsize=(25, 20))
_ = plot_tree(tree_pipe.steps[1][1],
              feature_names=X.columns,
              filled=True)
```



Iteration 2

This iteration attempts some initial pruning with values < 18 for `max_depth` and some higher values for `min_samples_leaf` and `min_samples_split`.

```
In [28]: tree_pipe2 = Pipeline([('ss', StandardScaler()),
                                ('clf', DecisionTreeClassifier(random_state=123))])

param_grid = {
    'clf__class_weight': ['balanced', {0:1, 1:200}],
    'clf__max_depth': [8, 16],
    'clf__min_samples_leaf': [1, 3, 6, 9],
```

```
'clf__min_samples_split': [2, 4, 6, 8]
}
```

```
In [29]: tree_grid2 = GridSearchCV(tree_pipe2, param_grid=param_grid, cv=3, scoring='reca
```

```
In [30]: tree_grid2.fit(X_train, y_train)
```

```
Out[30]: GridSearchCV(cv=3,
                      estimator=Pipeline(steps=[('ss', StandardScaler()),
                                                  ('clf',
                                                  DecisionTreeClassifier(random_state=12
3))]),
                      param_grid={'clf__class_weight': ['balanced', {0: 1, 1: 200}],
                                  'clf__max_depth': [8, 16],
                                  'clf__min_samples_leaf': [1, 3, 6, 9],
                                  'clf__min_samples_split': [2, 4, 6, 8]},
                      scoring='recall')
```

```
In [31]: tree_grid2.best_estimator_
```

```
Out[31]: Pipeline(steps=[('ss', StandardScaler()),
                          ('clf',
                           DecisionTreeClassifier(class_weight='balanced', max_depth=8,
                                                    min_samples_leaf=9,
                                                    random_state=123))])
```

```
In [32]: y_pred = tree_grid2.predict(X_test)
```

```
print('Train Report')
print(classification_report(y_train, tree_grid2.predict(X_train)))
print('\n')
print('Test Report')
print(classification_report(y_test, y_pred))
```


Train Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	0.95	1.00	0.97	1524
accuracy			1.00	691072
macro avg	0.97	1.00	0.99	691072
weighted avg	1.00	1.00	1.00	691072

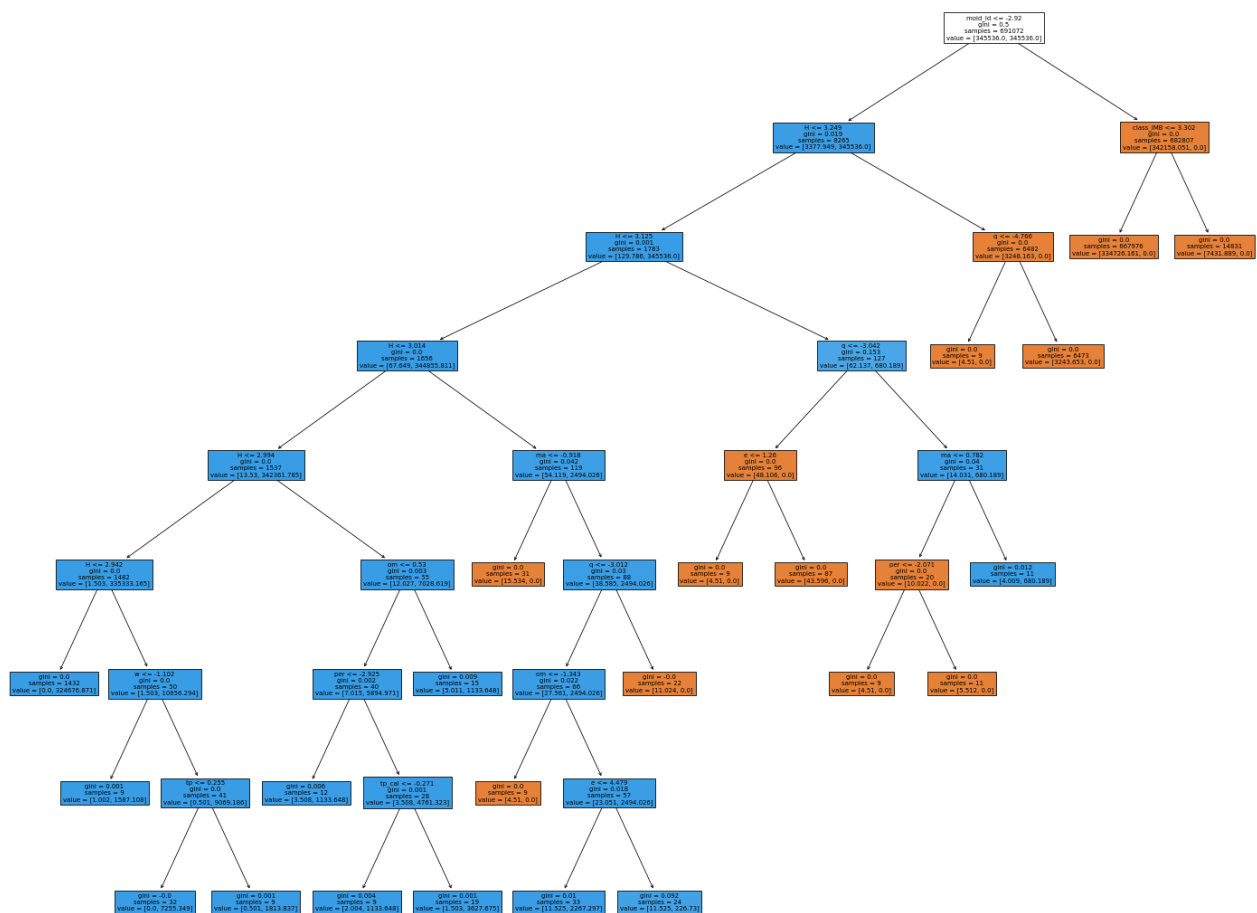
Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.94	0.99	0.96	541
accuracy			1.00	230358
macro avg	0.97	1.00	0.98	230358
weighted avg	1.00	1.00	1.00	230358

```
In [60]: plot_confusion_matrix(tree_grid2, X_test, y_test, cmap=plt.cm.Blues,
                                display_labels=['Non-Hazardous', 'Potentially Hazardous'],
```




```
In [61]: fig = plt.figure(figsize=(25, 20))
_ = plot_tree(tree_grid2.best_estimator_.steps[1][1],
               feature_names=X.columns,
               filled=True)
```



```
In [62]: plt.figure(figsize=(20, 15))
sns.barplot(x = tree_grid2.best_estimator_.steps[1][1].feature importances,
```

```
y = X.columns)
plt.title('Feature Importances');
```



`moid_ld` (closeness of orbit to orbit of Earth) is the most significant feature by far, followed by `H` (visual magnitude). Other features do not seem to have a signal in the feature importances plot, which is concerning.

```
In [36]: tree_grid2.best_params_
```

```
Out[36]: {'clf__class_weight': 'balanced',
          'clf__max_depth': 8,
          'clf__min_samples_leaf': 9,
          'clf__min_samples_split': 2}
```

The model chose the lower end of `max_depth`, the higher end of `min_samples_leaf`, and the default `min_samples_split`.

Iteration 3

This iteration searches around the best parameters of the previous iteration.

```
In [6]: tree_pipe3 = Pipeline([('ss', StandardScaler()),
                              ('clf', DecisionTreeClassifier(random_state=123))])

param_grid = {
    'clf__class_weight': ['balanced', {0:1, 1:200}],
    'clf__max_depth': [4, 6, 8, 12],
    'clf__min_samples_leaf': [7, 8, 9, 10, 11],
    'clf__min_samples_split': [2, 3, 4]
}
```

```
In [7]: tree_grid3 = GridSearchCV(tree_pipe3, param_grid=param_grid, cv=3, scoring='reca
```

```
In [8]: tree_grid3.fit(X_train, y_train)
```

```
Out[8]: GridSearchCV(cv=3,
                    estimator=Pipeline(steps=[('ss', StandardScaler()),
                                              ('clf',
                                               DecisionTreeClassifier(random_state=12
3))]),
                    param_grid={'clf__class_weight': ['balanced', {0: 1, 1: 200}],
                                'clf__max_depth': [4, 6, 8, 12],
                                'clf__min_samples_leaf': [7, 8, 9, 10, 11],
                                'clf__min_samples_split': [2, 3, 4]},
                    scoring='recall')
```

```
In [9]: y_pred = tree_grid3.predict(X_test)

print('Train Report')
print(classification_report(y_train, tree_grid3.predict(X_train)))
print('\n')
print('Test Report')
print(classification_report(y_test, y_pred))
```

Train Report


	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	0.90	1.00	0.95	1524
accuracy			1.00	691072
macro avg	0.95	1.00	0.97	691072
weighted avg	1.00	1.00	1.00	691072

Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.90	0.99	0.94	541
accuracy			1.00	230358
macro avg	0.95	1.00	0.97	230358
weighted avg	1.00	1.00	1.00	230358

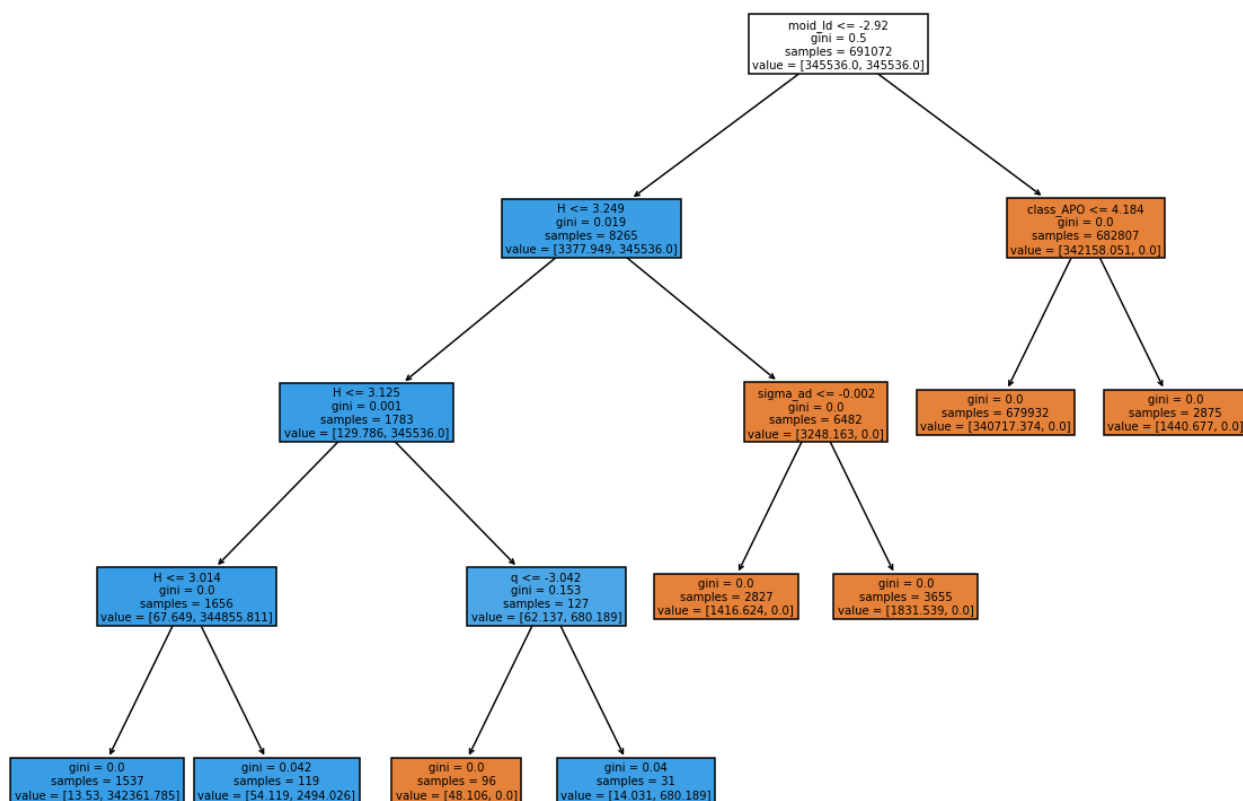
```
In [10]: sns.set_context('talk', font_scale=0.65)

plot_confusion_matrix(tree_grid3, X_test, y_test, cmap=plt.cm.Blues,
                      display_labels=['Non-Hazardous', 'Potentially Hazardous'],
plt.title('Decision Tree Confusion Matrix')
plt.savefig('Images/DT-matrix.png', bbox_inches='tight');
```

We have sacrificed some precision in order to identify more potentially hazardous asteroids correctly.

```
In [15]: fig = plt.figure(figsize=(20, 15))
_ = plot_tree(tree_grid3.best_estimator_.steps[1][1],
              feature_names=X.columns,
              filled=True)
plt.title('Decision Tree')
plt.savefig('Images/decision-tree.png');
```




```

In [18]: sns.set_context('talk')
sns.set_style('darkgrid')

plt.figure(figsize=(20, 15))
sns.barplot(x = tree_grid3.best_estimator_.steps[1][1].feature_importances_,
            y = X.columns, color='#682dd3')
plt.title('Feature Importances')
plt.xlabel('mean decrease in impurity')
plt.ylabel('feature')
plt.tight_layout()
plt.savefig('Images/DT-importances.png');

```



Only H and moid_ld have a signal on the feature importances plot.

```
In [43]: tree_grid3.best_params_
```

```
Out[43]: {'clf__class_weight': 'balanced',
          'clf__max_depth': 4,
          'clf__min_samples_leaf': 7,
          'clf__min_samples_split': 2}
```

Iteration 4

```
In [45]: tree_pipe4 = Pipeline([('ss', StandardScaler()),
                                ('clf', DecisionTreeClassifier(random_state=123,
                                                                min_samples_split=2))])

param_grid = {
    'clf__class_weight': ['balanced', {0:1, 1:200}],
    'clf__max_depth': [3, 4, 5],
    'clf__min_samples_leaf': [6, 7, 8]
}

tree_grid4 = GridSearchCV(tree_pipe4, param_grid=param_grid, cv=3, scoring='reca
```

```
In [46]: tree_grid4.fit(X_train, y_train)
```

```
Out[46]: GridSearchCV(cv=3,
                      estimator=Pipeline(steps=[('ss', StandardScaler()),
                                                  ('clf',
                                                  DecisionTreeClassifier(random_state=12
                                                  3))]),
```

```
param_grid={'clf__class_weight': ['balanced', {0: 1, 1: 200}],
            'clf__max_depth': [3, 4, 5],
            'clf__min_samples_leaf': [6, 7, 8]},
scoring='recall')
```

```
In [47]: y_pred = tree_grid4.predict(X_test)

print('Train Report')
print(classification_report(y_train, tree_grid4.predict(X_train)))
print('\n')
print('Test Report')
print(classification_report(y_test, y_pred))
```

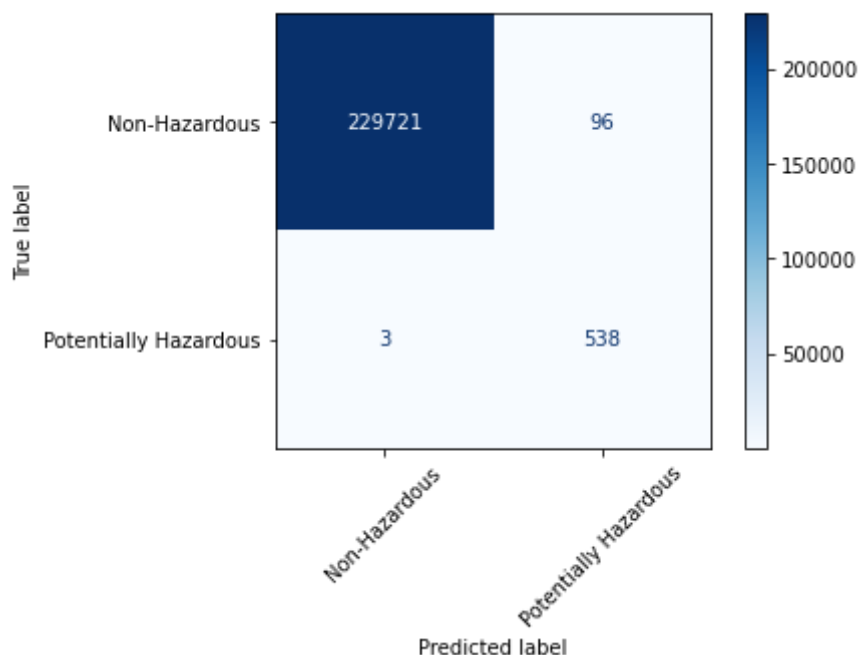
Train Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	0.85	1.00	0.92	1524
accuracy			1.00	691072
macro avg	0.93	1.00	0.96	691072
weighted avg	1.00	1.00	1.00	691072

Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.85	0.99	0.92	541
accuracy			1.00	230358
macro avg	0.92	1.00	0.96	230358
weighted avg	1.00	1.00	1.00	230358

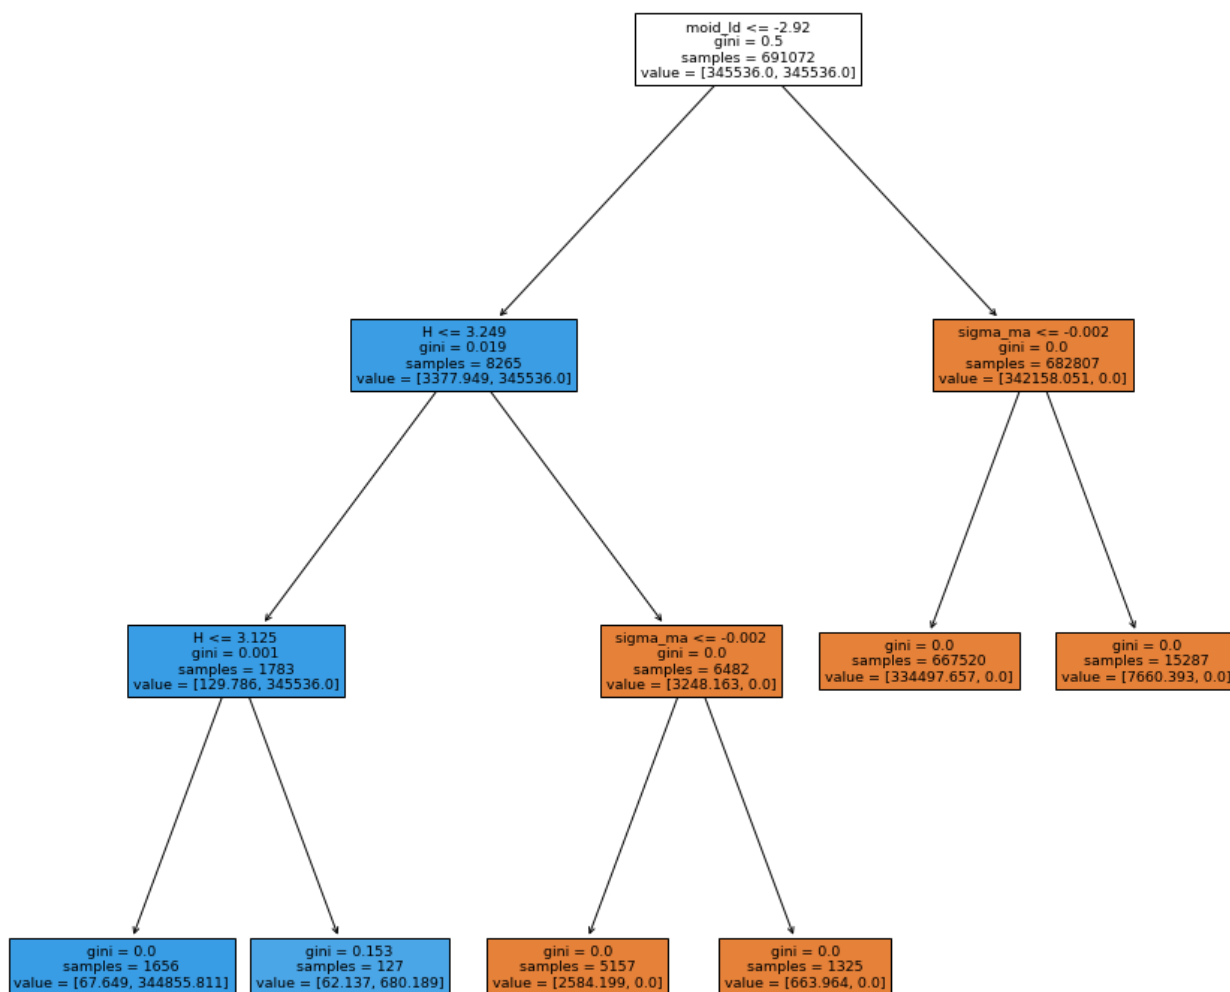
```
In [48]: plot_confusion_matrix(tree_grid4, X_test, y_test, cmap=plt.cm.Blues,
                               display_labels=['Non-Hazardous', 'Potentially Hazardous'],
```



Continuing to prioritize recall while gridsearching parameters does not seem to increase recall for the positive class (model has still only missed 3 potentially hazardous asteroids), but now

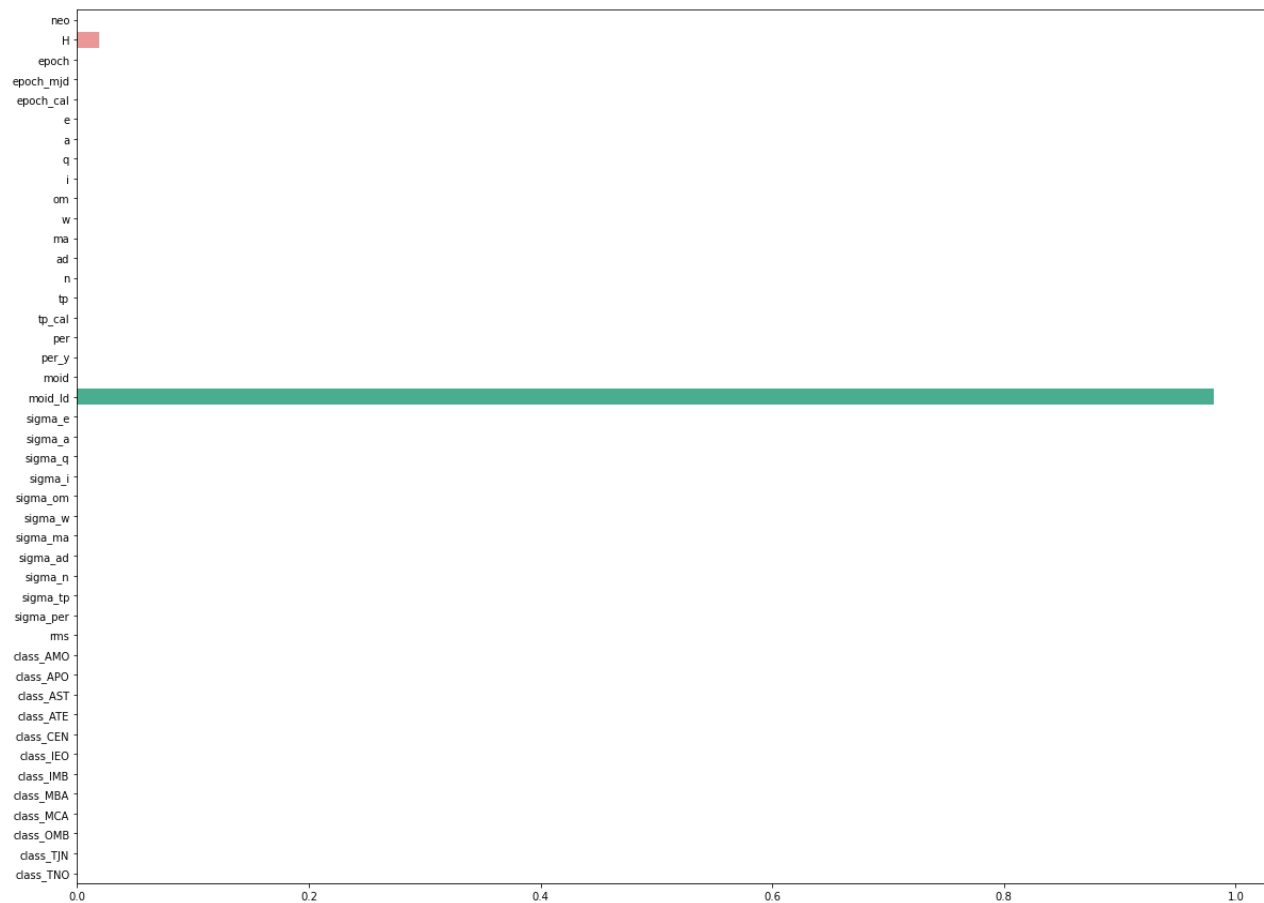
there are more false positives than the previous run (96 as opposed to 62, from the previous iteration).

```
In [49]: fig= plt.figure(figsize=(15, 15))
_ = plot_tree(tree_grid4.best_estimator_.steps[1][1],
              feature_names=X.columns,
              filled=True)
```



```
In [50]: plt.figure(figsize=(20, 15))
sns.barplot(x = tree_grid4.best_estimator_.steps[1][1].feature_importances_,
            y = X.columns)
```

Out[50]: <AxesSubplot:>



In [51]: `tree_grid4.best_params_`

Out[51]: `{'clf__class_weight': 'balanced',
'clf__max_depth': 3,
'clf__min_samples_leaf': 6}`