

Random Forest

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
```

```
In [3]: df = pd.read_csv('data/data_no_fliers.csv')
df.drop('Unnamed: 0', axis=1, inplace=True)
df.head()
```

```
Out[3]:
```

	neo	pha	H	epoch	epoch_mjd	epoch_cal	e	a	q	i
0	0	0	3.40	2458600.5	58600	20190427.0	0.076009	2.769165	2.558684	10.594067
1	0	0	4.20	2459000.5	59000	20200531.0	0.229972	2.773841	2.135935	34.832932
2	0	0	5.33	2459000.5	59000	20200531.0	0.256936	2.668285	1.982706	12.991043
3	0	0	3.00	2458600.5	58600	20190427.0	0.088721	2.361418	2.151909	7.141771
4	0	0	6.90	2459000.5	59000	20200531.0	0.190913	2.574037	2.082619	5.367427

5 rows × 45 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 921430 entries, 0 to 921429
Data columns (total 45 columns):
#   Column      Non-Null Count  Dtype
---  -
0   neo          921430 non-null  int64
1   pha          921430 non-null  int64
2   H            921430 non-null  float64
3   epoch        921430 non-null  float64
4   epoch_mjd    921430 non-null  int64
5   epoch_cal    921430 non-null  float64
6   e            921430 non-null  float64
7   a            921430 non-null  float64
8   q            921430 non-null  float64
9   i            921430 non-null  float64
10  om           921430 non-null  float64
11  w            921430 non-null  float64
12  ma           921430 non-null  float64
13  ad           921430 non-null  float64
14  n            921430 non-null  float64
15  tp           921430 non-null  float64
16  tp_cal       921430 non-null  float64
```

```

17 per          921430 non-null float64
18 per_y        921430 non-null float64
19 moid         921430 non-null float64
20 moid_ld      921430 non-null float64
21 sigma_e      921430 non-null float64
22 sigma_a      921430 non-null float64
23 sigma_q      921430 non-null float64
24 sigma_i      921430 non-null float64
25 sigma_om     921430 non-null float64
26 sigma_w      921430 non-null float64
27 sigma_ma     921430 non-null float64
28 sigma_ad     921430 non-null float64
29 sigma_n      921430 non-null float64
30 sigma_tp     921430 non-null float64
31 sigma_per    921430 non-null float64
32 rms          921430 non-null float64
33 class_AMO    921430 non-null int64
34 class_APO    921430 non-null int64
35 class_AST    921430 non-null int64
36 class_ATE    921430 non-null int64
37 class_CEN    921430 non-null int64
38 class_IEO    921430 non-null int64
39 class_IMB    921430 non-null int64
40 class_MBA    921430 non-null int64
41 class_MCA    921430 non-null int64
42 class_OMB    921430 non-null int64
43 class_TJN    921430 non-null int64
44 class_TNO    921430 non-null int64

```

dtypes: float64(30), int64(15)

memory usage: 316.3 MB

```

In [5]: X = df.drop('pha', axis=1)
        y = df['pha']

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random

```

```

In [6]: ss = StandardScaler()

        X_train_scaled = ss.fit_transform(X_train)
        X_test_scaled = ss.transform(X_test)

```

```

In [7]: RF = RandomForestClassifier(class_weight={0:1, 1:200})
        RF.fit(X_train_scaled, y_train)

```

Out[7]: RandomForestClassifier(class_weight={0: 1, 1: 200})

```

In [8]: y_pred = RF.predict(X_test_scaled)

        print('Train Report')
        print(classification_report(y_train, RF.predict(X_train_scaled)))
        print('\n')
        print('Test Report')
        print(classification_report(y_test, y_pred))

```

Train Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	1.00	1.00	1.00	1524
accuracy			1.00	691072
macro avg	1.00	1.00	1.00	691072


weighted avg	1.00	1.00	1.00	691072
--------------	------	------	------	--------

Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.98	0.96	0.97	541
accuracy			1.00	230358
macro avg	0.99	0.98	0.99	230358
weighted avg	1.00	1.00	1.00	230358

The baseline model is overfit to the training data, as we let the tree depths grow out.

```
In [9]: plot_confusion_matrix(RF, X_test_scaled, y_test, cmap=plt.cm.Blues,
                             display_labels=['Non-Hazardous', 'Potentially Hazardous'],
```



```
In [10]: depths = [estimator.tree_.max_depth for estimator in RF.estimators_]
          print('Mean depth:', np.mean(depths))
          depths
```

Mean depth: 26.93

```
Out[10]: [23,
          23,
          27,
          25,
          25,
          31,
          27,
          36,
          24,
          24,
          21,
          23,
          23,
          37,
          25,
          26,
```

24,
18,
26,
35,
24,
20,
26,
17,
30,
32,
29,
29,
25,
23,
28,
34,
21,
20,
35,
24,
43,
23,
29,
29,
32,
25,
25,
33,
35,
27,
30,
26,
22,
23,
24,
22,
22,
31,
22,
29,
28,
24,
22,
28,
27,
41,
24,
23,
23,
29,
23,
19,
32,
29,
38,
22,
25,
32,
27,
26,
24,
23,
21,
28,
26,

```

36,
34,
34,
24,
34,
24,
26,
34,
27,
21,
28,
19,
24,
23,
28,
41,
31,
25,
24]

```

```
In [11]: RF.get_params()
```

```

Out[11]: {'bootstrap': True,
'ccp_alpha': 0.0,
'class_weight': {0: 1, 1: 200},
'criterion': 'gini',
'max_depth': None,
'max_features': 'auto',
'max_leaf_nodes': None,
'max_samples': None,
'min_impurity_decrease': 0.0,
'min_impurity_split': None,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 100,
'n_jobs': None,
'oob_score': False,
'random_state': None,
'verbose': 0,
'warm_start': False}

```

Iteration 2

This iteration will search `max_depth` values less than 27 and a small range of values for `min_samples_leaf`. `min_samples_leaf` is kept at 2, the smallest possible value. This is also the value that worked best for the decision tree model, and one would expect the trees in a forest to be more grown out.

```

In [8]: RF2 = RandomForestClassifier(random_state=123)

param_grid = {
    'class_weight': ['balanced', {0:1, 1:200}],
    'n_estimators': [50, 100],
    'max_depth': [7, 14, 21],
    'min_samples_leaf': [1, 3, 6]
}

RF2_grid = GridSearchCV(RF2, param_grid=param_grid, cv=3, scoring='recall')

```

```
In [ ]: RF2_grid.fit(X_train_scaled, y_train)
```

```
In [15]: y_pred = RF2_grid.predict(X_test_scaled)

print('Train Report')
print(classification_report(y_train, RF2_grid.predict(X_train_scaled)))
print('\n')
print('Test Report')
print(classification_report(y_test, y_pred))
```

Train Report


	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	0.85	1.00	0.92	1524
accuracy			1.00	691072
macro avg	0.93	1.00	0.96	691072
weighted avg	1.00	1.00	1.00	691072

Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.85	1.00	0.92	541
accuracy			1.00	230358
macro avg	0.92	1.00	0.96	230358
weighted avg	1.00	1.00	1.00	230358

The model is no longer overfit to the training data.

```
In [39]: plot_confusion_matrix(RF2_grid, X_test_scaled, y_test, cmap=plt.cm.Blues,
                                display_labels=['Non-Hazardous', 'Potentially Hazardous'],
```



```
In [17]: RF2_grid.best_params_
```

```
Out[17]: {'class_weight': {0: 1, 1: 200},
           'max_depth': 7,
```

```
'min_samples_leaf': 3,
'n_estimators': 100}
```

Iteration 3

Searching around `max_depth = 7` and `min_samples_leaf = 3`

```
In [10]: RF3 = RandomForestClassifier(random_state=123)

class_weights = ['balanced', {0:1, 1:200}]

param_grid = {
    'class_weight': class_weights,
    'max_depth': [6, 7, 8, 9],
    'min_samples_leaf': [2, 3, 4, 5]
}

RF3_grid = GridSearchCV(RF3, param_grid=param_grid, cv=3, scoring='recall')
```

```
In [11]: RF3_grid.fit(X_train_scaled, y_train)
```

```
Out[11]: GridSearchCV(cv=3, estimator=RandomForestClassifier(random_state=123),
    param_grid={'class_weight': ['balanced', {0: 1, 1: 200}],
                'max_depth': [6, 7, 8, 9],
                'min_samples_leaf': [2, 3, 4, 5]},
    scoring='recall')
```

```
In [12]: y_pred = RF3_grid.predict(X_test_scaled)

print('Train Report')
print(classification_report(y_train, RF3_grid.predict(X_train_scaled)))
print('\n')
print('Test Report')
print(classification_report(y_test, y_pred))
```

Train Report


	precision	recall	f1-score	support
0	1.00	1.00	1.00	689548
1	0.50	1.00	0.67	1524
accuracy			1.00	691072
macro avg	0.75	1.00	0.83	691072
weighted avg	1.00	1.00	1.00	691072

Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.52	1.00	0.69	541
accuracy			1.00	230358
macro avg	0.76	1.00	0.84	230358
weighted avg	1.00	1.00	1.00	230358

```
In [13]: sns.set_context('paper')

plot_confusion_matrix(RF3_grid, X_test_scaled, y_test, cmap=plt.cm.Blues,
    display_labels=['Non-Hazardous', 'Potentially Hazardous'],
```



```
In [14]: RF3_grid.best_params_
```

```
Out[14]: {'class_weight': 'balanced', 'max_depth': 6, 'min_samples_leaf': 2}
```

Iteration 4

Last gridsearch will omit `max_depth = 6` to see if we can address the high number of false positives.

```
In [17]: RF4 = RandomForestClassifier(random_state=123)

param_grid = {
    'class_weight': class_weights,
    'max_depth': [7, 8, 9],
    'min_samples_leaf': [2, 3, 4, 5]
}

RF4_grid = GridSearchCV(RF4, param_grid=param_grid, cv=3, scoring='recall')
```

```
In [18]: RF4_grid.fit(X_train_scaled, y_train)
```

```
Out[18]: GridSearchCV(cv=3, estimator=RandomForestClassifier(random_state=123),
    param_grid={'class_weight': ['balanced', {0: 1, 1: 200}],
    'max_depth': [7, 8, 9],
    'min_samples_leaf': [2, 3, 4, 5]},
    scoring='recall')
```

```
In [19]: y_pred = RF4_grid.predict(X_test_scaled)

print('Train Report')
print(classification_report(y_train, RF4_grid.predict(X_train_scaled)))
print('\n')
print('Test Report')
print(classification_report(y_test, y_pred))
```

Train Report


precision recall f1-score support

	Random-Forest			
0	1.00	1.00	1.00	689548
1	0.85	1.00	0.92	1524
accuracy			1.00	691072
macro avg	0.93	1.00	0.96	691072
weighted avg	1.00	1.00	1.00	691072

Test Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	229817
1	0.85	1.00	0.92	541
accuracy			1.00	230358
macro avg	0.92	1.00	0.96	230358
weighted avg	1.00	1.00	1.00	230358

```
In [30]: sns.set_context('talk', font_scale=0.65)
plot_confusion_matrix(RF4_grid, X_test_scaled, y_test, cmap=plt.cm.Blues,
                      display_labels=['Non-Hazardous', 'Potentially Hazardous'],
plt.title('Random Forest Confusion Matrix')
plt.savefig('Images/RF-matrix.png', bbox_inches='tight')
plt.savefig('Images/RF-matrix-hr.png', dpi=200, bbox_inches='tight');
```



This model missed the same number of phas as the previous iteration, but there are much fewer false positives.

```
In [21]: RF4_grid.best_params_
```


```
Out[21]: {'class_weight': {0: 1, 1: 200}, 'max_depth': 7, 'min_samples_leaf': 3}
```

```
In [22]: RF4 = RandomForestClassifier(random_state=123, class_weight={0:1, 1:200}, max_de
RF4.fit(X_train_scaled, y_train)
```

```
Out[22]: RandomForestClassifier(class_weight={0: 1, 1: 200}, max_depth=7,
                                min_samples_leaf=3, random_state=123)
```

```
In [33]: sns.set_context('talk')
sns.set_style('darkgrid')

plt.figure(figsize=(20, 15))
sns.barplot(x = RF4.feature_importances_,
            y = X.columns, color='#682dd3')
plt.title('Feature Importances')
plt.xlabel('mean decrease in impurity')
plt.ylabel('feature')
plt.tight_layout()
plt.savefig('Images/RF-importances.png')
plt.savefig('Images/RF-importances.png', dpi=200);
```



```
In [34]: from sklearn.decomposition import PCA

pca = PCA(n_components=5)
X_transformed = pca.fit_transform(X_train_scaled)
```

```
In [35]: pca.explained_variance_ratio_
```

```
Out[35]: array([0.19685156, 0.19090429, 0.10615482, 0.06094324, 0.04144515])
```

```
In [36]: pca.components_
```

```
Out[36]: array([[ 1.27542039e-02,  1.52473281e-02, -1.58175961e-02,
                  -1.58175961e-02, -1.58290957e-02,  1.03213098e-02,
```

```
-1.84319288e-02, -1.91784300e-02, -9.85767667e-04,  
-1.53721468e-04, -2.08587819e-04, -7.20260799e-04,  
-1.27459463e-02, 1.72449064e-02, -1.40138695e-02,  
-1.39574146e-02, -1.80149625e-02, -1.80149625e-02,  
-1.89824635e-02, -1.89824635e-02, 2.66314755e-01,  
3.25919764e-01, 2.61916406e-01, 2.95404804e-01,  
2.18694776e-01, 3.27418165e-01, 3.14062198e-01,  
3.25233226e-01, 3.10534332e-01, 3.15362133e-01,  
3.27373260e-01, 6.31337904e-05, 6.32400184e-03,  
9.99037782e-03, -4.70669720e-04, 4.97793807e-03,  
-1.61738770e-03, 6.51942931e-04, 4.60691810e-03,  
-7.75711353e-03, 4.49679407e-03, -4.95202787e-03,  
-3.48840881e-03, -0.00000000e+00],  
[-2.08029712e-01, -2.40662799e-01, 7.68968374e-02,  
7.68968374e-02, 7.71649506e-02, -1.60202421e-01,  
3.22067675e-01, 3.23862930e-01, 8.51264474e-03,  
-3.11736182e-03, -6.85094289e-04, 1.00094776e-02,  
2.32694313e-01, -2.97992289e-01, 6.94856364e-02,  
6.92863096e-02, 3.15567275e-01, 3.15567275e-01,  
3.20495242e-01, 3.20495242e-01, 1.68380010e-02,  
2.03905489e-02, 1.65722501e-02, 1.84570066e-02,  
1.38336019e-02, 2.06088716e-02, 1.97079832e-02,  
2.03446266e-02, 1.93577886e-02, 1.98349027e-02,  
2.05025611e-02, -5.31778459e-04, -1.03243575e-01,  
-1.60801526e-01, 1.00234944e-02, -8.67192111e-02,  
3.13270355e-02, -1.12712827e-02, -8.50278890e-02,  
1.22477613e-01, -7.17198709e-02, 9.29215879e-02,  
6.48945204e-02, -0.00000000e+00],  
[ 5.04435854e-02, 5.83809628e-02, -4.32708461e-01,  
-4.32708461e-01, -4.32535015e-01, 6.25047189e-02,  
1.04064623e-01, 4.68136216e-02, 3.48130986e-02,  
1.15992828e-03, -1.14031619e-05, -6.77016015e-02,  
1.26494873e-01, -6.27556860e-02, -4.15360787e-01,  
-4.13288498e-01, 1.09148295e-01, 1.09148295e-01,  
5.20714451e-02, 5.20714451e-02, -5.84144254e-03,  
-7.50950291e-03, -5.97478597e-03, -5.68749618e-03,  
-4.52441514e-03, -8.34013461e-03, -8.37113078e-03,  
-7.24631210e-03, -7.69435015e-03, -8.23011758e-03,  
-7.37636539e-03, -4.35061294e-03, 2.99983142e-02,  
4.96721615e-02, 1.25854102e-02, -1.83330542e-02,  
3.15973158e-02, -1.80021235e-03, -3.87973934e-02,  
-6.10775270e-02, 2.17590275e-02, 6.40394092e-02,  
4.46874827e-02, -0.00000000e+00],  
[ 3.49079774e-01, 1.17435399e-01, 9.78174596e-02,  
9.78174596e-02, 9.74774021e-02, 3.84050341e-01,  
1.69901726e-01, -7.41751051e-02, 2.10554277e-01,  
2.57886182e-03, 1.99757534e-03, 3.50296037e-02,  
3.40140764e-01, -5.92203625e-03, 1.15746698e-01,  
1.15089382e-01, 1.98510619e-01, 1.98510619e-01,  
-4.59436561e-02, -4.59436561e-02, 6.20690275e-04,  
9.58514072e-04, 1.94201484e-04, -2.13043390e-04,  
2.90036706e-03, 1.89919267e-03, 1.57317827e-03,  
1.08053548e-03, 6.13289381e-04, 1.97692477e-03,  
1.07268351e-03, 1.95833442e-03, 2.32637610e-01,  
2.62978844e-01, 4.46999929e-02, 3.40105736e-02,  
1.22431685e-01, 4.91122677e-03, -3.78098086e-02,  
-3.84291767e-01, 1.18280885e-01, 2.50424547e-01,  
1.24844518e-01, -0.00000000e+00],  
[-4.81640185e-02, -4.83080825e-02, -1.86687861e-02,  
-1.86687861e-02, -1.82572273e-02, -3.41812612e-01,  
-6.22154540e-02, 1.43063750e-01, 2.93117501e-01,  
1.13414211e-02, -2.13854392e-02, 3.83502783e-02,  
-2.27383270e-01, 1.68970365e-01, -3.23543596e-03,  
-2.24895674e-03, -4.07083816e-02, -4.07083816e-02,  
1.57011240e-01, 1.57011240e-01, -4.27329590e-03,
```

```

1.55537734e-03, -4.25554016e-03, -2.61907550e-03,
-2.22158113e-03, 4.13474069e-04, 2.42752590e-03,
1.36051718e-03, 2.48219699e-03, 1.63276697e-03,
1.34093337e-03, 1.26008409e-02, -4.90320037e-02,
-7.18510381e-02, -3.56658417e-02, 1.24289327e-01,
-4.29049908e-02, 4.11308753e-02, 5.18329579e-01,
-4.67414008e-01, 1.10942505e-01, 3.19623044e-01,
5.12157241e-02, -0.00000000e+00]])

```

```
In [37]: pd.DataFrame(data=pca.components_, columns=X.columns).iloc[1]
```

```

Out[37]: neo          -0.208030
H             -0.240663
epoch         0.076897
epoch_mjd     0.076897
epoch_cal     0.077165
e            -0.160202
a             0.322068
q             0.323863
i             0.008513
om           -0.003117
w           -0.000685
ma            0.010009
ad            0.232694
n           -0.297992
tp            0.069486
tp_cal       0.069286
per           0.315567
per_y        0.315567
moid         0.320495
moid_ld      0.320495
sigma_e      0.016838
sigma_a      0.020391
sigma_q      0.016572
sigma_i      0.018457
sigma_om     0.013834
sigma_w      0.020609
sigma_ma     0.019708
sigma_ad     0.020345
sigma_n      0.019358
sigma_tp     0.019835
sigma_per    0.020503
rms         -0.000532
class_AMO   -0.103244
class_APO   -0.160802
class_AST    0.010023
class_ATE   -0.086719
class_CEN    0.031327
class_IEO   -0.011271
class_IMB   -0.085028
class_MBA    0.122478
class_MCA   -0.071720
class_OMB    0.092922
class_TJN    0.064895
class_TNO   -0.000000
Name: 1, dtype: float64

```

```
In [ ]:
```